

Computação Paralela

Mest. Int. Engenharia Computacional

Ano letivo 2019/2020

Manuel Barroso, Nuno Lau

Superscalar Compiler and Processor

Sequential
source code


```
a = foo(b);  
for (i=0, i<
```

Superscalar compiler

Find independent
operations

Schedule
operations

Sequential
machine code



Superscalar processor

Check instruction
dependencies

Schedule
execution

- Extensions to the base ISA that allow a form of vector computation
- “Vectors” are implemented in dedicated registers
 - MMX: 64 bits
 - SSE: 128 bits
 - AVX: 256 bits
 - AVX-512: 512 bits
- Registers of N bits may be used as vectors of $2 \times (N/2)$ elements, $4 \times (N/4)$ elements, etc.
- One multimedia instruction applies simultaneously to all elements of a register


- MMX
 - 57 instructions added to Pentium
- Extended MMX
- SSE
 - Pentium III
 - 71 instructions (52 FP SIMD, 19 MMX)
- SSE2
 - 144 new instructions (Pentium 4)
 - 128 bit registers
 - Cache-control
- SSE3
 - 13 new instructions (Pentium 4 - 2004)
 - Horizontal processing of values in a register
- SSSE3
 - 32 new instructions (Core)
- SSE4
 - 54 instructions (Penryn - 2008)

- 3D Now!
 - AMD
 - 45 instructions (21 FP SIMD, 19 MMX, 5 DSP)
- AltiVec
 - Motorola
 - 162 instructions
- AVX
 - Sandy Bridge – 2011
 - 256 bit registers
 - Instructions with 3 operands
- AVX2
 - Haswell New Instructions, 2013
 - Broadcast/permute operations on data elements
 - Vector shift instructions with variable-shift count per data element
 - Instructions to fetch non-contiguous data elements from memory
- AVX-512
 - Proposed in 2013, First Processor 2016
 - 512 bit registers
 - Several extensions: Foundation, Prefetch, Vector Neural Network, etc.


- 8 registers of 64 bits
 - MM0, MM1, ..., MM7
 - Are implemented on the same hardware as the FP registers: ST0, ..., ST7
 - This allowed to maintain compatibility with existing operating systems

MMX data types


- 8 bytes array



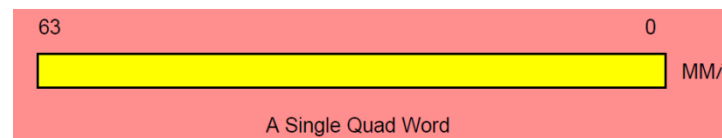
- 4 words array (16 bit/word)



- 2 double words array (32 bits)

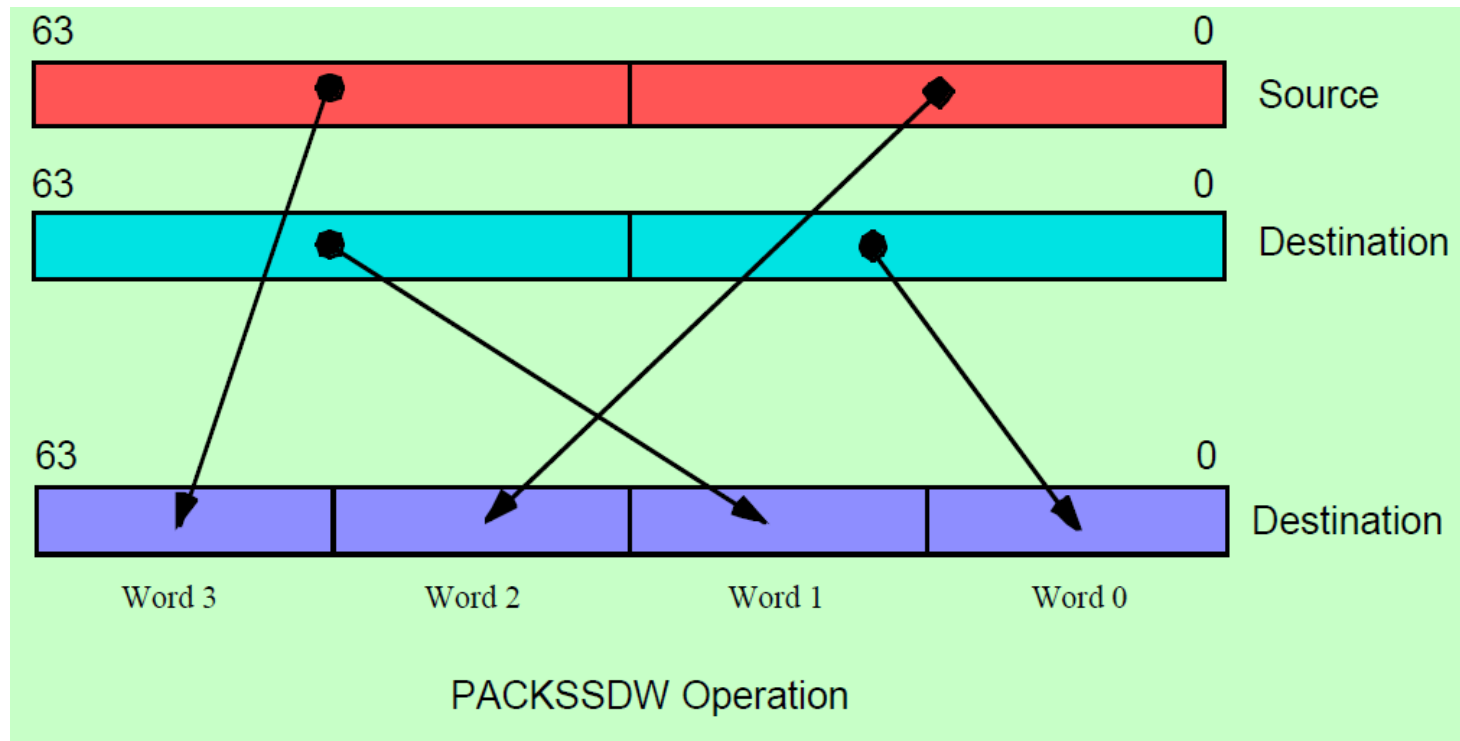


- Quadword (64 bits)

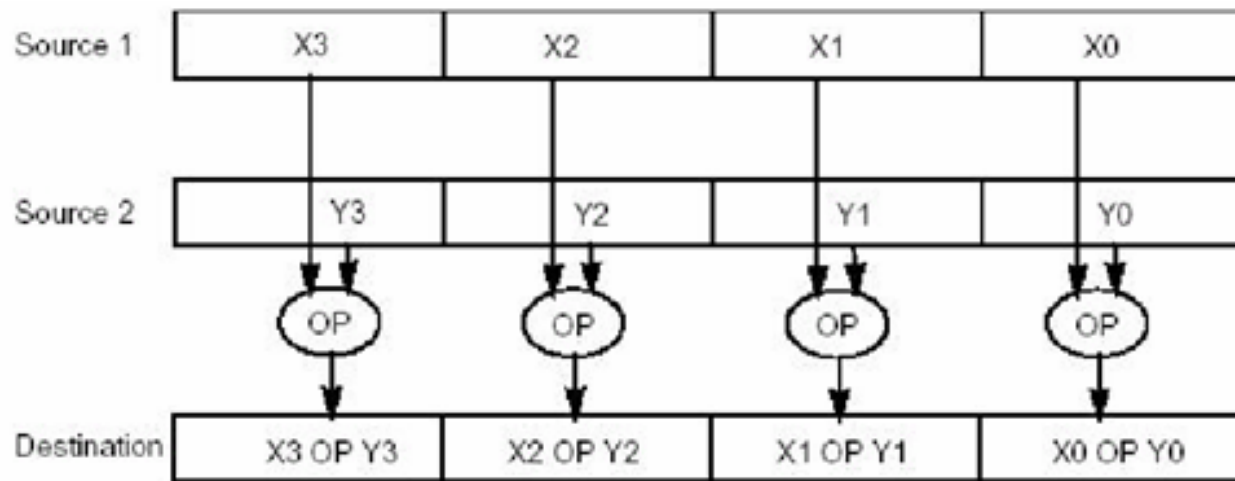


- Data transfer
 - Move data between registers, memory and MMX registers
 - `movd` (32 bits) and `movq` (64 bits)
- Conversion
 - Converts larger data types to smaller data types and vice-versa
- Packed arithmetic
- Comparisons
- Logic operations
- Shift and Rotate
- EMMS
 - Prepares the processor to execute FP code again

- Conversion



- Packed arithmetic
 - SIMD: *Single Instruction Multiple Data*
 - paddb (8bit), paddw (16bit), padd (32bit)
 - paddb, paddsb (*signed saturation*), paddusb (*unsigned saturation*)



MMX instructions

Packed Arithmetic	Wrap Around	Signed Sat	Unsigned Sat
Addition	PADD	PADDs	PADDUS
Subtraction	PSUB	PSUBs	PSUBUS
Multiplication	PMULL/H		
Multiply & add	PMADD		
Shift right Arithmetic	PSRA		
Compare	PCMPcc		
Conversions	Regular	Signed Sat	Unsigned Sat
Pack		PACKSS	PACKUS
Unpack	PUNPCKL/H		
Logical Operations	Packed	Full 64-bit	
And		PAND	
And not		PANDN	
Or		POR	
Exclusive or		PXOR	
Shift left	PSLL	PSLL	
Shift right	PSRL	PSRL	
Transfers and Memory Operations	32-bit	64-bit	
Register-register move	MOVD	MOVQ	
Load from memory	MOVD	MOVQ	
Store to memory	MOVD	MOVQ	
Miscellaneous			
Empty multimedia state	EMMS		

- 8 registers of 128 bits
 - XMM0, ..., XMM7
- 4 FPs simple precision (SSE)
- 2 FPs double precision (SSE2)
- 16 bytes (SSE2)
- 8 words (SSE2)
- 4 double words (SSE2)
- 1 inteiro 128bit (SSE2)

Code:

```
int A[size], B[size], C[size];
```

```
...
```

```
for (i = 0 ; i < size ; i++)
```

```
    C[i] = A[i] + B[i];
```

Parallelism?

Sum arrays

```
movdqa (%eax,%edx,4), %xmm0    # load A[i] to A[i+3]
movdqa (%ebx,%edx,4), %xmm1    # load B[i] to B[i+3]
padd    %xmm0, %xmm1           # CCCC = AAAA + BBBB
movdqa %xmm1, (%ecx,%edx,4)     # store C[i] to C[i+3]
addl    $4, %edx               # i += 4
```

- `movdqa (%eax,%edx,4), %xmm0`
 - **mov**: transferência
 - **dq**: double quad
 - **a**: align
 - **(%eax,%edx,4)**: $\text{eax} + 4 * \text{edx}$

Sum elements of an array

Code:

```
int A[size], total;
```

```
...
```

```
for (i = 0 ; i < size ; i++)  
    total += A[i];
```

Parallelism?

Sum elements of an array

Restructured code:

```
int A[size], temp[4], total;
temp[0]=temp[1]=temp[2]=temp[3]=0;
for (i = 0 ; i < size ; i+=4) {
    temp[0] += A[i];    temp[1] += A[i+1];
    temp[2] += A[i+2]; temp[3] += A[i+3];
}
total = temp[0]+temp[1]+temp[2]+temp[3];
```

Parallelism?

- Internal product of 2 vectors (**a** and **b**)

```
length6 = (size/24)*24
for(; i < length6; i += 24){
    __asm__ volatile
    (// instruction comment
      "\n\t movdqa 0x00(%0),%%xmm2 \t#" "\n\t movdqa 0x10(%0),%%xmm3 \t#"
      "\n\t movdqa 0x20(%0),%%xmm4 \t#" "\n\t movdqa 0x30(%0),%%xmm5 \t#"
      "\n\t movdqa 0x40(%0),%%xmm6 \t#" "\n\t movdqa 0x50(%0),%%xmm7 \t#"

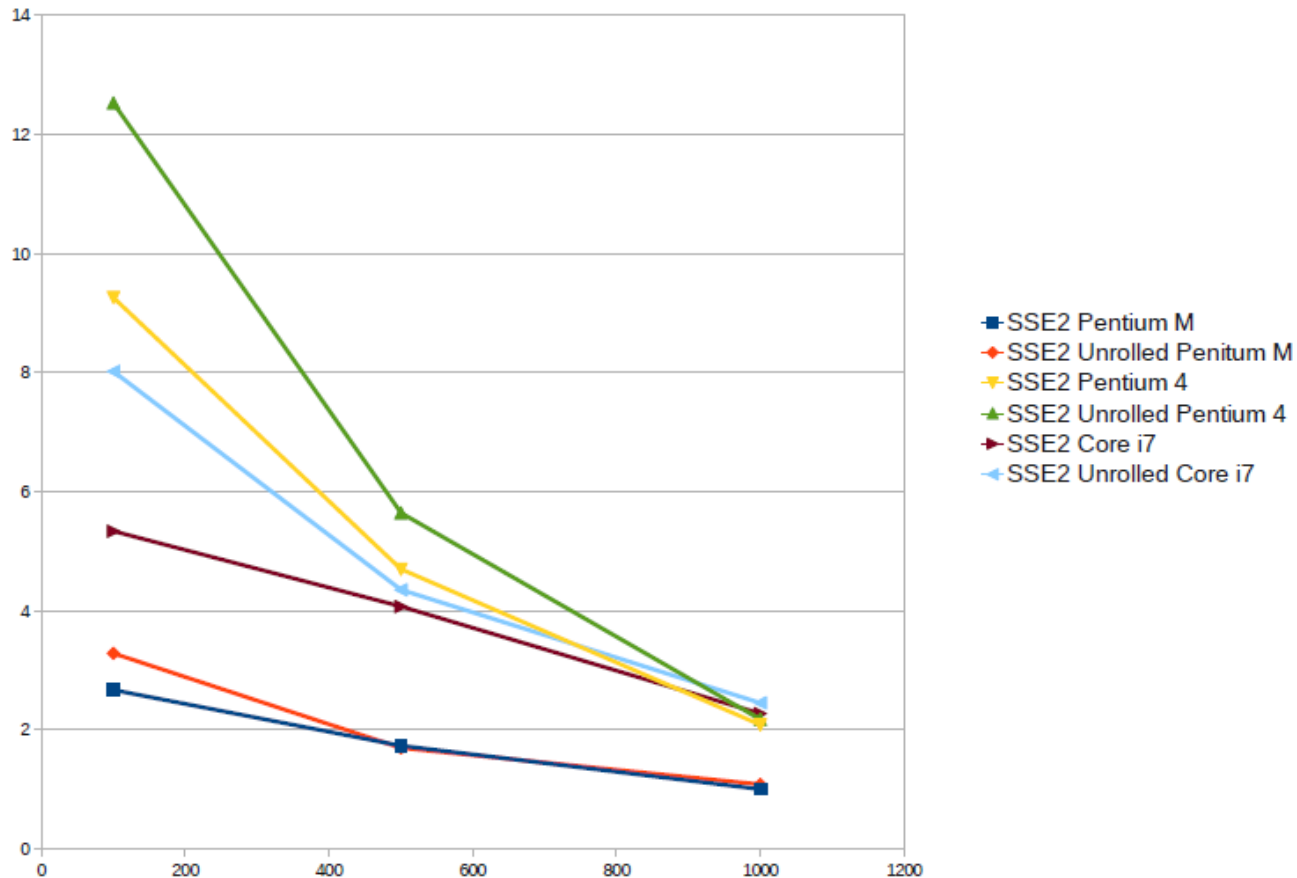
      "\n\t mulps 0x00(%1),%%xmm2 \t#" "\n\t mulps 0x10(%1),%%xmm3 \t#"
      "\n\t mulps 0x20(%1),%%xmm4 \t#" "\n\t mulps 0x30(%1),%%xmm5 \t#"
      "\n\t mulps 0x40(%1),%%xmm6 \t#" "\n\t mulps 0x50(%1),%%xmm7 \t#"

      "\n\t addps %%xmm2,%%xmm0 \t#" "\n\t addps %%xmm3,%%xmm1 \t#"
      "\n\t addps %%xmm4,%%xmm0 \t#" "\n\t addps %%xmm5,%%xmm0 \t#"
      "\n\t addps %%xmm6,%%xmm1 \t#" "\n\t addps %%xmm7,%%xmm0 \t#"

      :
      : "r" (a+i), // %0
      : "r" (b+i) // %1
    );
}
```

SSE with loop unrolling

- Effect of loop unrolling in internal product



- Comparison with other libraries

