

# Computação Paralela

Mest. Int. Engenharia Computacional

Ano letivo 2019/2020

Manuel Barroso, Nuno Lau

- **Open specifications for Multi Processing** *via collaborative work between interested parties from the hardware and software industry, government and academia*

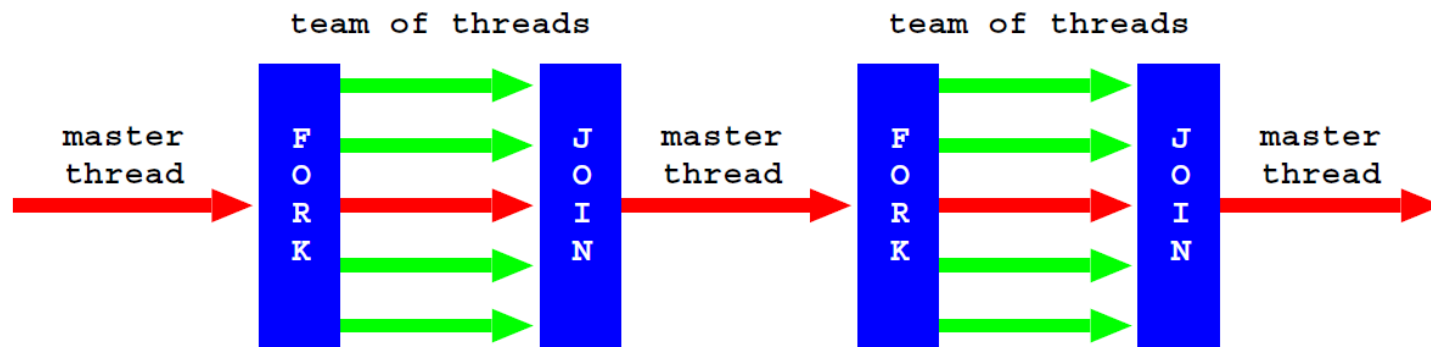
- **Shared Memory Programming Model**
- Cooperation of several hardware and software companies (Sun, Intel, Fujitsu, IBM, AMD, HP, SGI, ...)
- Parallel Programming API for multiprocessor / multicore architectures
- Languages: C/C++ or Fortran
- OS: Unix/Linux or Windows
- **OpenMP is a specification not an implementation!**

- Portable programming model for shared memory architectures
- Simple set of programming directives
- Enable incremental parallelism for sequential programs
- Efficient implementations for different problems

- Explicit parallelism
  - Programmer is responsible for annotating execution tasks and synchronization points
  - Embedded compiler directives
- Implicit multithreading
  - Process is a set of threads that communicate through shared variables
  - Creation and termination of threads is performed implicitly by the execution environment
    - Global address space is shared by all threads
    - Variables may be shared or private for each thread

# OpenMP fork-join execution

- Program initiates with a single (master) thread
- Executes sequentially until parallel region defined by OpenMP constructor, and then:
  - Master thread forks “team of threads”
  - Parallel region code is executed concurrently by all threads (including master thread)
  - There is an implicit barrier at the end of parallel region
  - “team of threads” terminates and master thread continues sequentially



From CP@FCUP

- API includes
  - Compiler directives: `#pragma omp <directive>`
  - Library of functions declared at `omp.h`
  - Environment variables (`OMP_NUM_THREADS; ...`)
- OpenMP code
  - Must include `omp.h` header file  
`#include <omp.h>`
  - To compile with gcc use `-fopenmp` option  
`gcc -fopenmp myprog.c ...`

```
#pragma omp parallel [clause, ...]
```

- Defines parallel region
- Number of threads determined by:
  - Only master if `if(expr)` clause is used and `expr` is false
  - Number defined by `num_threads(expr)` clause
  - Number defined by last call to `omp_set_num_threads(expr)`
  - Number defined by environment variable `OMP_NUM_THREADS`
  - Implementation dependent



```
#include <stdio.h>
#include <omp.h>

int main ()
{
    #pragma omp parallel
    {
        int ID = omp_get_thread_num();

        printf("Hello (%d) ", ID);
        printf("World (%d) \n", ID);
    }
}
```

- **`int omp_get_thread_num(void) ;`**
  - Returns thread id from 0 to n-1, where n is the number of threads
  - Master thread has id 0
- **`int omp_get_num_threads(void) ;`**
  - Returns the number of active threads
- **`void omp_set_num_threads(int num_threads) ;`**
  - Specifies the maximum number of threads to be used in the next parallel regions
  - Can only be called from sequential code

# Shared and private variables

```
#pragma omp parallel shared(list1) private(list2)
```

- Shared variables
  - Shared by all threads
  - Avoid race conditions!
- Private variables
  - Duplicated in each thread
  - Initial value is undefined
  - Value after parallel region is undefined
- Variables are shared by default
  - Can be changed by using **default(mod)** clause
    - mod: **private**, **shared**, **none**
- Variables declared inside parallel region are private

```
#pragma omp parallel firstprivate(list)
```

- Firstprivate variables
  - Duplicated in each thread
  - Initial value is copied from value of variable with the same name
  - Can be used to improve efficiency
    - Latency of reading private variable may be lower than latency of reading shared variable by all threads

# Example: Dot product

```
int dot_product(int* u, int* v, int n) {  
    int r = 0;  
    for (int i = 0; i < n; i++) {  
        r += u[i] * v[i];  
    }  
    return r;  
}
```

- Partition iterations / vector among threads
- Reduction of partial results

# Example: Dot product

First version (incorrect!):

```
int dot_product(int* u, int* v, int n) {  
    int r = 0;  
    #pragma omp parallel  
    {  
        int rank = omp_get_thread_num();  
        int my_n = n / omp_get_num_threads();  
        for (int i = rank * my_n; i < (rank+1)*my_n; i++)  
            r += u[i] * v[i];  
    }  
    return r;  
}
```

- Manual partition
- Error prone

# Example: Dot product

Better version:

```
int dot_product(int* u, int* v, int n) {  
    int r = 0;  
    #pragma omp parallel  
    {  
        #pragma omp for reduction(+:r)  
        for (int i = 0; i < n; i++)  
            r += u[i] * v[i];  
    }  
    return r;  
}
```

- **#pragma omp for**
  - distributes iterations among threads
- Reduction is automatically performed by using **reduction** clause
- Scheduling may be controlled through **schedule** clause

# Example: Dot product

Even better version:

```
int dot_product(int* u, int* v, int n) {  
    int r = 0;  
    #pragma omp parallel for reduction(+:r)  
    for (int i = 0; i < n; i++) {  
        r += u[i] * v[i];  
    }  
    return r;  
}
```

- **#pragma omp parallel for**
  - creates parallel region and distributes iterations among threads



- OpenMP reduction clause:
  - `reduction (op : list)`
- Inside a parallel or a work-sharing construct:
  - A local copy of each list variable is made and initialized depending on the “op” (e.g. 0 for “+”).
  - Updates occur on the local copy.
  - Local copies are reduced into a single value and combined with the original global value.
- The variables in “list” must be shared in the enclosing parallel region.

```
double ave=0.0, A[MAX]; int i;
#pragma omp parallel for reduction (+:ave)
for (i=0;i< MAX; i++) {
    ave + = A[i];
}
ave = ave/MAX;
```

# OpenMP Reduction

Operator	Initial value
+	0
*	1
-	0
min	Largest pos. number
max	Most neg. number

C/C++ only

Operator	Initial value
&	1
	0
^	0
&&	1
	0

- OpenMP has constructs to enable mutual exclusion among threads
- 3 different ways:
  - **omp critical**
    - Defines critical region
    - Can be used for any code
  - **omp atomic**
    - Mutual exclusion for atomic variable updates
    - Only for simple memory updates
      - $x \text{ binop} = \text{expr}$ ,  $x++$ ;  $++x$ ;  $x--$ ;  $--x$
  - Explicit use of locks

- **omp critical**

```
int dot_product(int* u, int* v, int n) {  
    int r = 0;  
    #pragma omp parallel  
    {  
        int private_r = 0;  
        #pragma omp for  
        for (int i = 0; i < n; i++)  
            private_r += u[i] * v[i];  
        #pragma omp critical  
        r += private_r;  
    }  
    return r;  
}
```

- **omp atomic**

```
int dot_product(int* u, int* v, int n) {  
    int r = 0;  
    #pragma omp parallel  
    {  
        int private_r = 0;  
        #pragma omp for  
        for (int i = 0; i < n; i++)  
            private_r += u[i] * v[i];  
        #pragma omp atomic  
        r += private_r;  
    }  
    return r;  
}
```

- Locks

```
int dot_product(int* u, int* v, int n) {  
    int r = 0;  
    omp_lock_t lock;  
    omp_init_lock(&lock); // initialize lock  
    #pragma omp parallel  
    {  
        int private_r = 0;  
        #pragma omp for  
        for (int i = 0; i < n; i++)  
            private_r += u[i] * v[i];  
        omp_set_lock(&lock); // acquire lock  
        r += private_r;  
        omp_unset_lock(&lock); // release lock  
    }  
    omp_destroy_lock(&lock); // destroy lock  
    return r;  
}
```

- **omp master**
  - Executed only by master thread
  - Ignored by others; no synchronization
- **omp single**
  - Executed by a single thread
  - Implicit barrier for all threads
    - Except `nowait` clause is used
- **omp barrier**
  - Explicit barrier

- **schedule (static [ , chunk ] )**
  - Each thread has fixed number of iterations
- **schedule (dynamic [ , chunk ] )**
  - Each thread grabs “chunk” iterations off a queue until all iterations have been handled
- **schedule (guided [ , chunk ] )**
  - Threads dynamically grab blocks of iterations
  - The size of the block starts large and shrinks down to size “chunk” as the calculation proceeds.
- **schedule (runtime)**
  - Schedule and chunk size taken from the OMP\_SCHEDULE environment variable (or the runtime library).
- **schedule (auto)**
  - Schedule is left up to the runtime to choose (does not have to be any of the above).