

Computação Paralela

Mest. Int. Engenharia Computacional

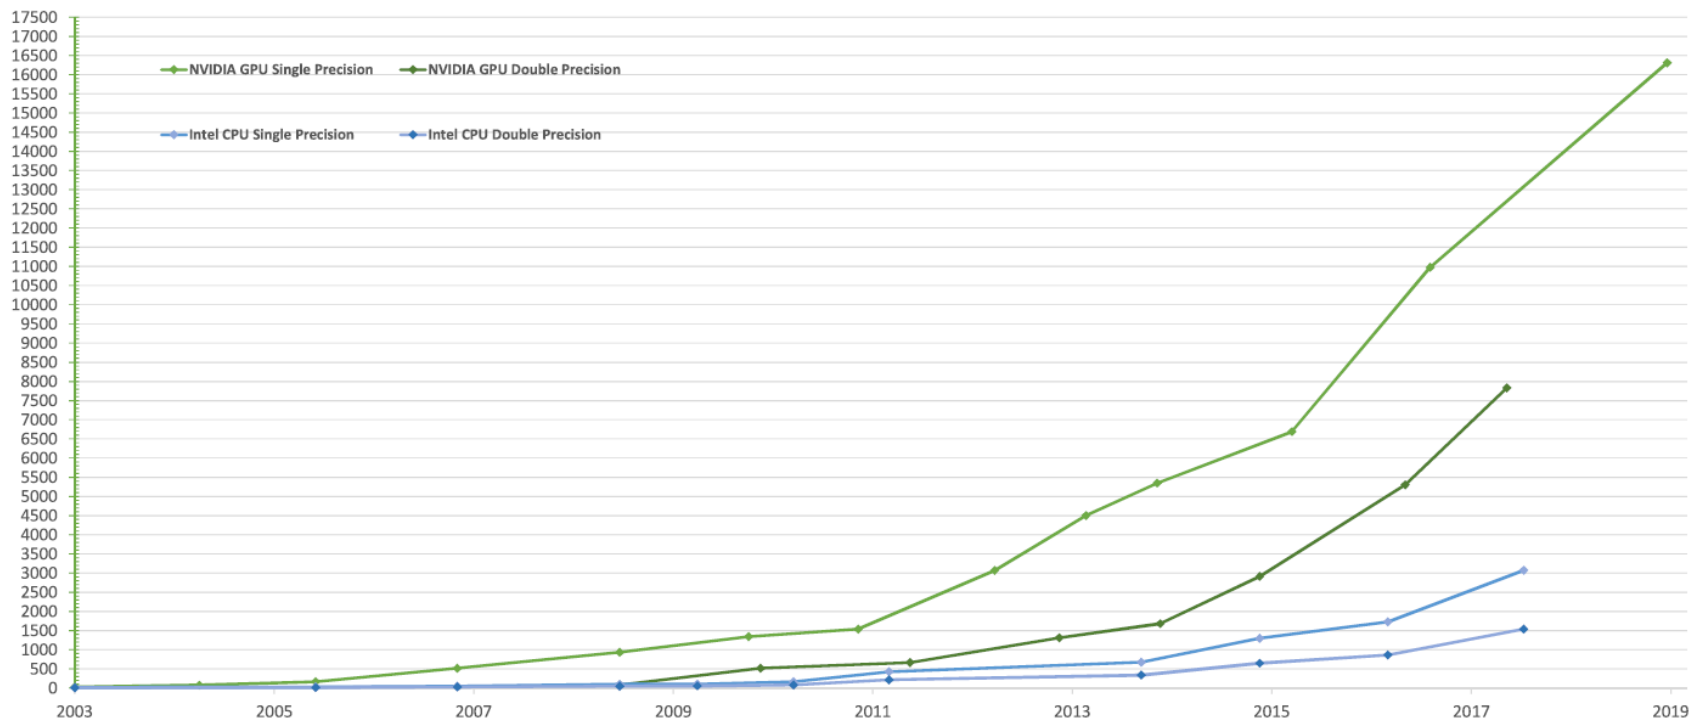
Ano letivo 2019/2020

Manuel Barroso, Nuno Lau


- Parallel general purpose computation model
- Introduced by NVIDIA in 2006
- Allows using the GPU for the execution of general purpose applications

Acknowledgement: This lecture is based on “CUDA C++ Programming Guide”, v10.2, NVIDIA


Theoretical GFLOP/s




Theoretical GB/s



CUDA – CPU and GPU Architecture

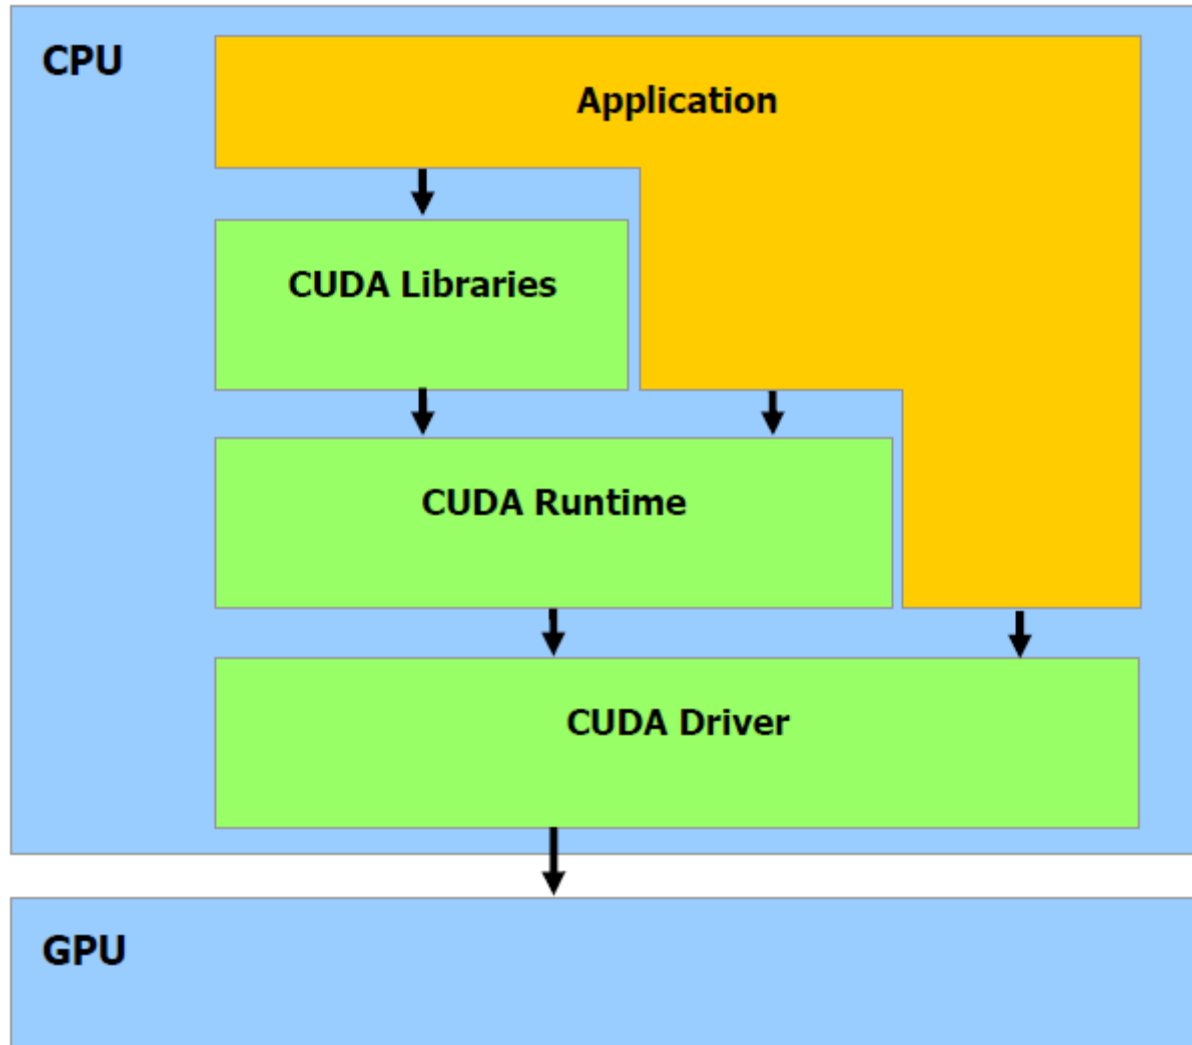


CPU



GPU

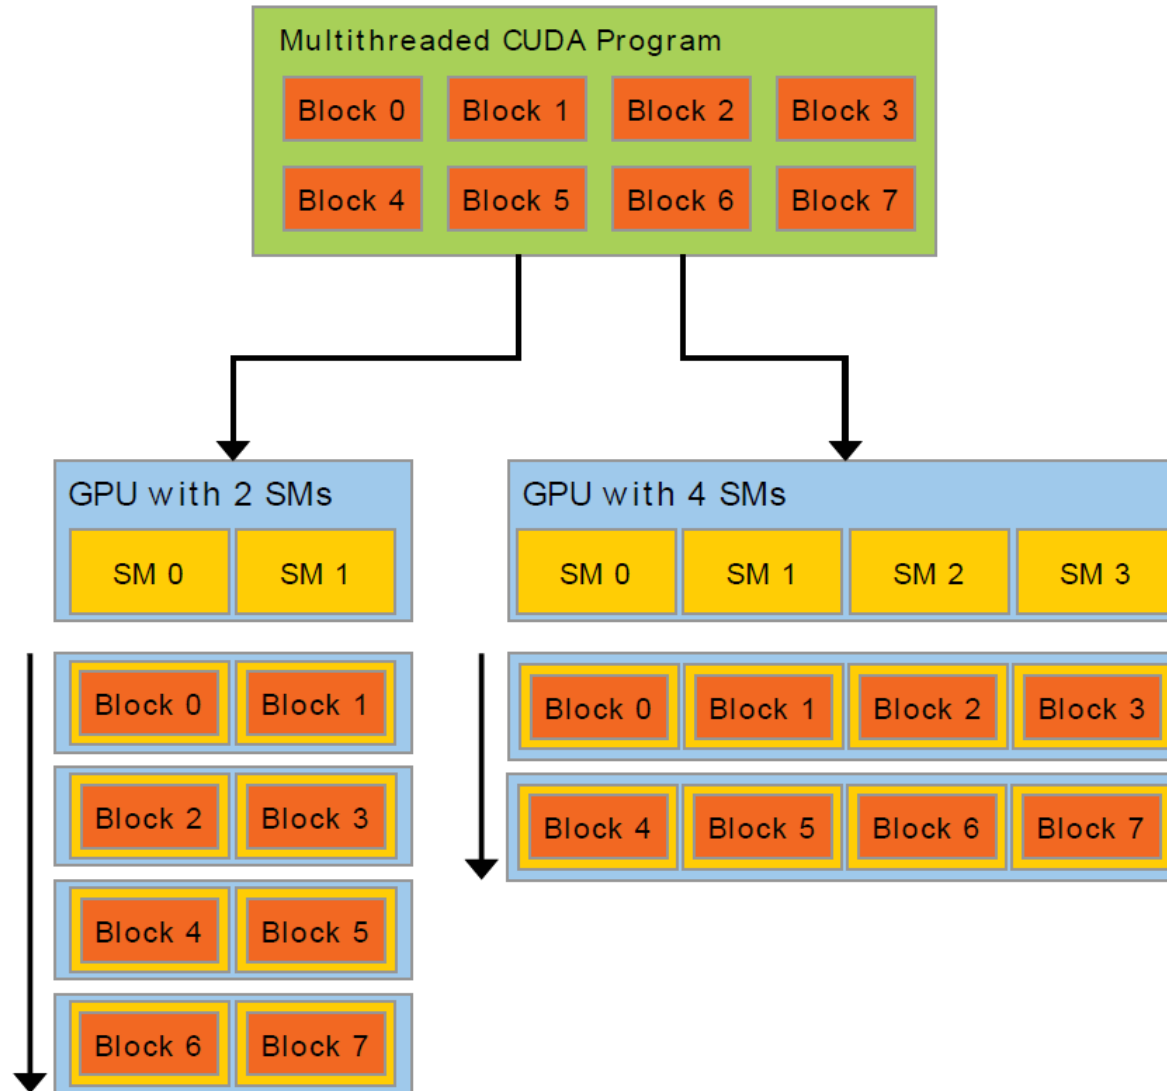
CUDA Architecture




CUDA Computing Applications

GPU Computing Applications						
Libraries and Middleware						
cuDNN TensorRT	cuFFT, cuBLAS, cuRAND, cuSPARSE	CULA MAGMA	Thrust NPP	VSIPPL, SVM, OpenCurrent	PhysX, OptiX, iRay	MATLAB Mathematica
Programming Languages						
C	C++	Fortran	Java, Python, Wrappers	DirectCompute	Directives (e.g., OpenACC)	
CUDA-enabled NVIDIA GPUs						
Turing Architecture (Compute capabilities 7.x)	DRIVE/JETSON AGX Xavier	GeForce 2000 Series		Quadro RTX Series	Tesla T Series	
Volta Architecture (Compute capabilities 7.x)	DRIVE/JETSON AGX Xavier				Tesla V Series	
Pascal Architecture (Compute capabilities 6.x)	Tegra X2	GeForce 1000 Series		Quadro P Series	Tesla P Series	
Maxwell Architecture (Compute capabilities 5.x)	Tegra X1	GeForce 900 Series		Quadro M Series	Tesla M Series	
Kepler Architecture (Compute capabilities 3.x)	Tegra K1	GeForce 700 Series GeForce 600 Series		Quadro K Series	Tesla K Series	
	EMBEDDED	CONSUMER DESKTOP, LAPTOP		PROFESSIONAL WORKSTATION	DATA CENTER	

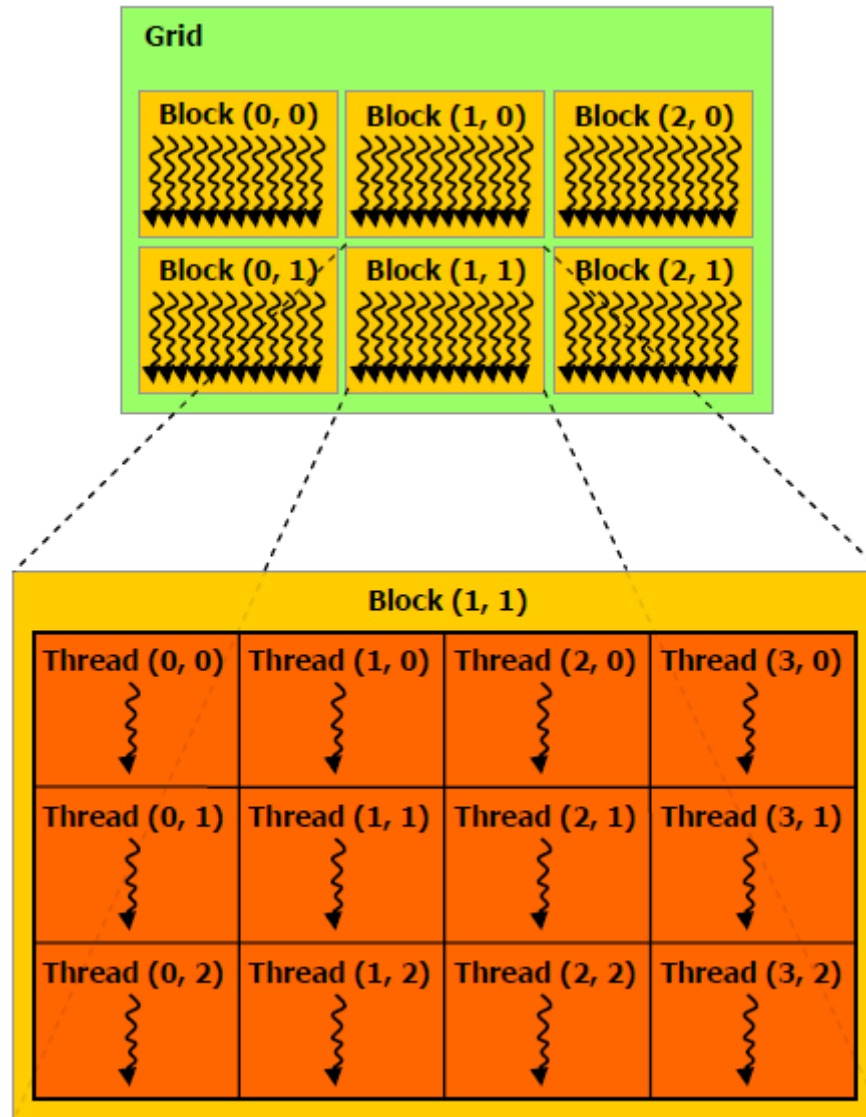
CUDA Automatic Scalability



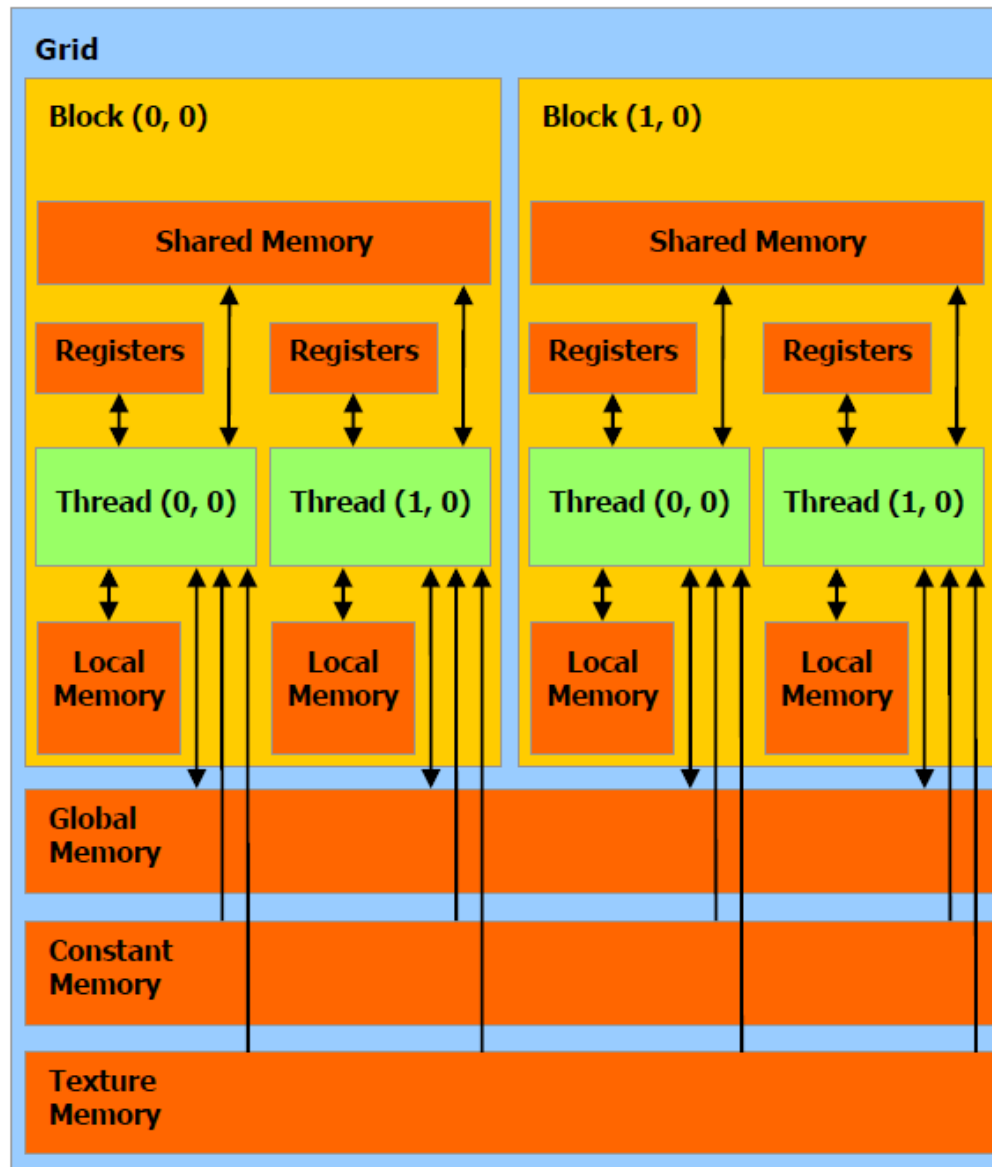
CUDA Kernel, Grid, Block, Thread



CUDA Grid of Threads



CUDA memory




CUDA memory

Thread

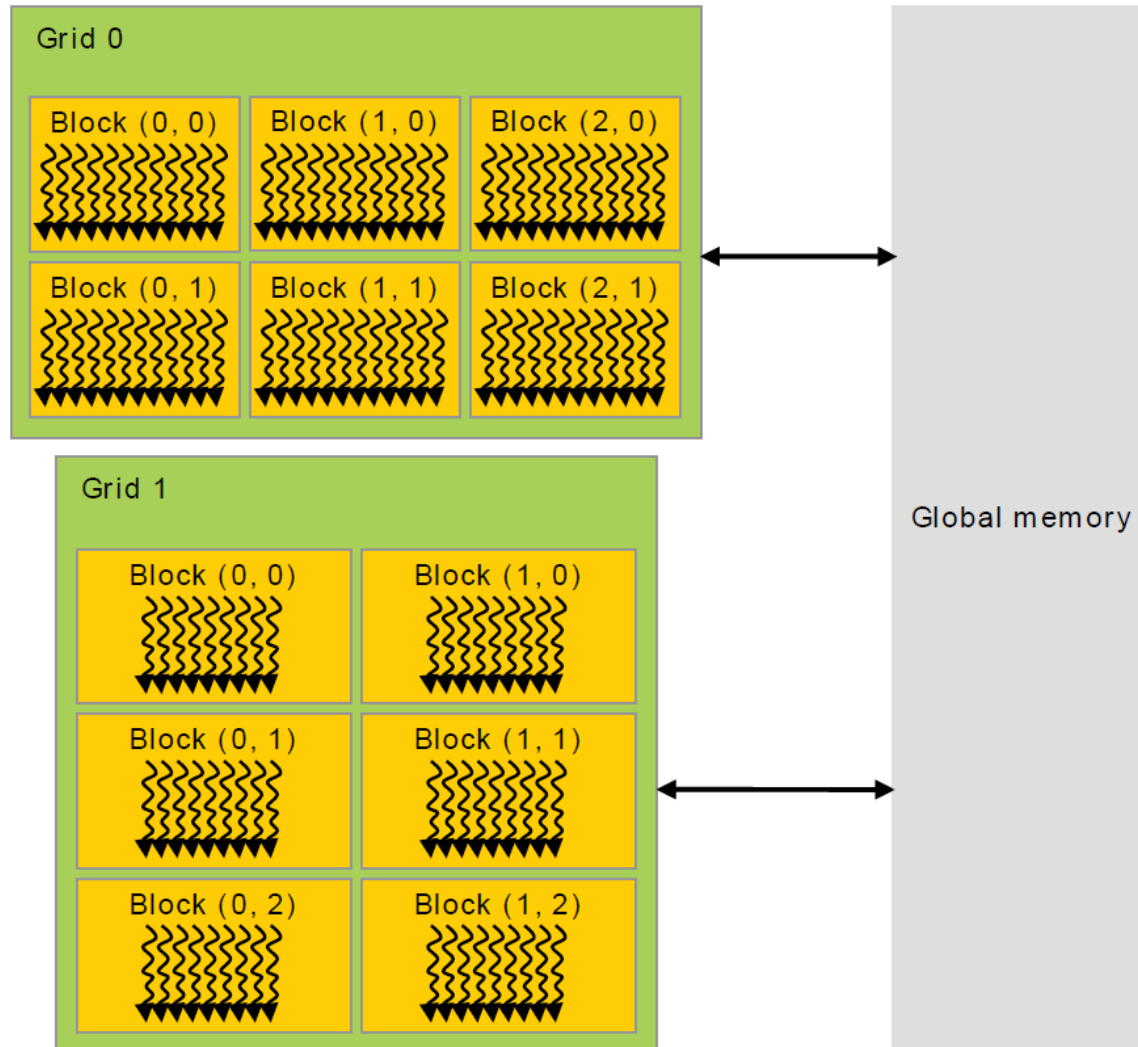


Per-thread local
memory



Per-block shared
memory

CUDA memory



Compute Capability 1.x

- Global memory (read and write)
 - Slow and uncached
 - Requires sequential and aligned 16 byte reads and writes to be fast
- Texture memory (read only)
 - Cache optimized for 2D spatial access pattern
- Constant memory
 - This is where constants and kernel arguments are stored
 - Slow, but with small cache
- Shared memory (16 kb per MP)
 - Fast, but take care of bank conflicts
 - Can be used to exchange data between threads in a block
- Local memory (used for data that does not fit into registers)
 - Slow and uncached
- Registers
 - Fastest, scope is thread local

Compute Capability 2.x

- Global memory (read and write)
 - Slow, but now with cache
- Texture memory (read only)
 - Cache optimized for 2D spatial access pattern
- Constant memory
 - Slow, but with cache (8 kb)
- Shared memory
 - Fast, but slightly different rules for bank conflicts now
- Local memory
 - Slow, but now with cache
- Registers (32768 32-bit registers per MP)

Compute Capability 3.x

- Global memory (read and write)
 - Generally cached in L2, but not in L1
- Constant memory
 - Cache shared by all functional units
- Shared memory
 - 32 banks with two addressing modes
- Local memory
 - Cached in L1

Compute Capability 5.x

- Global memory
 - Identical to cp 3.x
- Shared memory
 - 32 banks organized such that successive 32-bit words map to successive banks

CUDA Heterogeneous Programming

C Program
Sequential
Execution

Serial code

Parallel kernel
Kernel0<<<>>>()

Serial code

Parallel kernel
Kernel1<<<>>>()

Host

Device

Grid 0

Block (0, 0)

Block (1, 0)

Block (2, 0)

Block (0, 1)

Block (1, 1)

Block (2, 1)

Host

Device

Grid 1

Block (0, 0)

Block (1, 0)

Block (0, 1)

Block (1, 1)

Block (0, 2)

Block (1, 2)

```
// Kernel definition
__global__ void VecAdd(float* A, float* B, float* C)
{
    int i = threadIdx.x;
    C[i] = A[i] + B[i];
}

int main()
{
    ...
    // Kernel invocation with N threads
    VecAdd<<<1, N>>>(A, B, C);
}
```

- Kernel invocation

kernel<<< Dg, Db, Ns, S >>>

- **Dg** is the grid dimension (dim3 type)
- **Db** is the block dimension (dim3 type)
- **Ns** is the number of shared memory bytes
- **S** is the stream

- Block
 - `threadIdx.x`, `threadIdx.y`, `threadIdx.z`
 - Identify the thread in the block
 - `blockDim.x`, `blockDim.y`, `blockDim.z`
 - $\text{threadID} = x + y * \text{Dx} + z * \text{Dx} * \text{Dy}$
 - Threads are executed in *warps*
 - 32 threads of the same block with consecutive ids
 - Blocks should have more than 32 threads

- Grid
 - blockIdx.x, blockIdx.y, blockIdx.z
 - Identify the block in the grid
 - blockDim.x, blockDim.y, blockDim.z